Manejo de datos con R

Oscar Perpiñán Lamigueiro
http://oscarperpinan.github.io

setwd, getwd, dir

En `setwd` hay que especificar el directorio que contiene el repositorio.

```
getwd()
old <- setwd("~/github/intro")
dir()
```

```
dir(pattern='.R')
```

```
 [1] "birds.R"               "ClasesMetodos.R"       "datos.R"
 [4] "estadistica.R"         "factorDateCharacter.R" "Funciones.R"
 [7] "graficos.R"            "intro.R"               "raster.R"
[10] "zoo.R"
```

```
dir('data')
```

```
 [1] "aranjuez.csv"         "aranjuez.RData"       "bird_tracking.csv"
 [4] "CO2_GNI_BM.csv"       "El.Arenosillo.txt"    "eric.csv"
 [7] "InformeDatos.zip"     "nico.csv"             "NREL-Hawaii.csv"
[10] "radiacion_datas.csv" "sanne.csv"            "SIAR.csv"
[13] "SISmm2008_CMSAF.zip"
```

Lectura de datos con `read.table` o `read.csv`

- Función Genérica

```
dats <- read.table('data/aranjuez.csv', sep=',', header=TRUE)

head(dats)
```

```
           X TempAvg TempMax TempMin HumidAvg HumidMax WindAvg WindMax Rain
1 2004-01-01   4.044   10.71  -1.969     88.3     95.9   0.746   3.528    0
2 2004-01-02   5.777   11.52   1.247     83.3     98.5   1.078   6.880    0
3 2004-01-03   5.850   13.32   0.377     75.0     94.4   0.979   6.576    0
4 2004-01-04   4.408   15.59  -2.576     82.0     97.0   0.633   3.704    0
5 2004-01-05   3.081   14.58  -2.974     83.2     97.0   0.389   2.244    0
6 2004-01-06   2.304   11.83  -3.379     84.5     96.5   0.436   2.136    0
  Radiation        ET
1     5.490 0.5352688
2     6.537 0.7710499
3     8.810 0.8361229
4     9.790 0.6861381
5    10.300 0.5152422
6     9.940 0.4886631
```

- Función específica

```
aranjuez <- read.csv('data/aranjuez.csv')

head(aranjuez)
```

```
           X TempAvg TempMax TempMin HumidAvg HumidMax WindAvg WindMax Rain
1 2004-01-01   4.044   10.71  -1.969     88.3     95.9   0.746   3.528    0
2 2004-01-02   5.777   11.52   1.247     83.3     98.5   1.078   6.880    0
```

Inspeccionamos el resultado

```
names(aranjuez)
```

```
[1] "X"         "TempAvg"   "TempMax"   "TempMin"   "HumidAvg"  "HumidMax"
[7] "WindAvg"   "WindMax"   "Rain"      "Radiation" "ET"
```

```
head(aranjuez)
```

```
           X TempAvg TempMax TempMin HumidAvg HumidMax WindAvg WindMax Rain
1 2004-01-01   4.044   10.71  -1.969     88.3     95.9   0.746   3.528    0
2 2004-01-02   5.777   11.52   1.247     83.3     98.5   1.078   6.880    0
3 2004-01-03   5.850   13.32   0.377     75.0     94.4   0.979   6.576    0
4 2004-01-04   4.408   15.59  -2.576     82.0     97.0   0.633   3.704    0
5 2004-01-05   3.081   14.58  -2.974     83.2     97.0   0.389   2.244    0
6 2004-01-06   2.304   11.83  -3.379     84.5     96.5   0.436   2.136    0
  Radiation        ET
1     5.490 0.5352688
2     6.537 0.7710499
3     8.810 0.8361229
4     9.790 0.6861381
5    10.300 0.5152422
6     9.940 0.4886631
```

```
tail(aranjuez)
```

```
              X TempAvg TempMax TempMin HumidAvg HumidMax WindAvg WindMax  Rain
2893 2011-12-26   3.366   13.88  -3.397     81.5      100   0.556   3.263 0.000
2894 2011-12-27   2.222   13.33  -4.005     87.0      100   0.369   1.842 0.000
2895 2011-12-28   1.810   12.33  -4.682     85.0      100   0.540   3.401 0.203
2896 2011-12-29   2.512   11.92  -4.682     77.2      100   0.546   4.420 0.203
2897 2011-12-30   1.006   11.05  -5.822     79.7      100   0.446   2.832 0.000
2898 2011-12-31   2.263   12.67  -3.938     80.3      100   0.270   1.950 0.000
     Radiation        ET
2893      9.44 0.5358751
```

Inspeccionamos el resultado

```
summary(aranjuez)
```

```
         X          TempAvg          TempMax          TempMin
2004-01-01:   1   Min.   :-5.309   Min.   :-2.362   Min.   :-12.980
2004-01-02:   1   1st Qu.: 7.692   1st Qu.:14.530   1st Qu.:  1.515
2004-01-03:   1   Median :13.810   Median :21.670   Median :  7.170
2004-01-04:   1   Mean   :14.405   Mean   :22.531   Mean   :  6.888
2004-01-05:   1   3rd Qu.:21.615   3rd Qu.:30.875   3rd Qu.: 12.590
2004-01-06:   1   Max.   :30.680   Max.   :41.910   Max.   : 22.710
(Other)   :2892                                    NA's   :4
   HumidAvg         HumidMax         WindAvg          WindMax
Min.   : 19.89   Min.   : 35.88   Min.   :0.251   Min.   : 0.000
1st Qu.: 47.04   1st Qu.: 81.60   1st Qu.:0.667   1st Qu.: 3.783
Median : 62.58   Median : 90.90   Median :0.920   Median : 5.027
Mean   : 62.16   Mean   : 87.22   Mean   :1.174   Mean   : 5.208
3rd Qu.: 77.38   3rd Qu.: 94.90   3rd Qu.:1.431   3rd Qu.: 6.537
Max.   :100.00   Max.   :100.00   Max.   :8.260   Max.   :10.000
                 NA's   :13       NA's   :8       NA's   :128
     Rain          Radiation           ET
Min.   : 0.000   Min.   : 0.277   Min.   :0.000
1st Qu.: 0.000   1st Qu.: 9.370   1st Qu.:1.168
Median : 0.000   Median :16.660   Median :2.758
Mean   : 1.094   Mean   :16.742   Mean   :3.091
3rd Qu.: 0.200   3rd Qu.:24.650   3rd Qu.:4.926
Max.   :49.730   Max.   :32.740   Max.   :8.564
NA's   :4        NA's   :13       NA's   :18
```

Valores ausentes

- `NA` está definido como `logical`

```
class(NA)
```

```
[1] "logical"
```

- Operar con `NA` siempre produce un `NA`

```
1 + NA
```

```
[1] NA
```

- Esto es un «problema» al usar funciones

```
mean(aranjuez$Radiation)
```

```
[1] NA
```

```
mean(aranjuez$Radiation, na.rm = TRUE)
```

```
[1] 16.74176
```

Valores ausentes

Las funciones `is.na` y `anyNA` los identifican

```
anyNA(aranjuez)
```

```
[1] TRUE
```

```
which(is.na(aranjuez$Radiation))
```

```
[1] 1861 1867 1873 1896 1897 1908 1923 2153 2413 2587 2600 2603 2684
```

```
sum(is.na(aranjuez$Radiation))
```

```
[1] 13
```

Fechas

```
names(aranjuez)[1] <- "Date"

aranjuez$Date <- as.Date(aranjuez$Date)
```

```
class(aranjuez$Date)

summary(aranjuez$Date)
```

```
[1] "Date"
        Min.      1st Qu.       Median         Mean      3rd Qu.         Max.
"2004-01-01" "2005-12-29" "2008-01-09" "2008-01-03" "2010-01-02" "2011-12-31"
```

Fechas

- Podemos extraer información de un objeto `Date` con la función `format`[1]:

```
aranjuez$month <- as.numeric(
    format(aranjuez$Date, '%m'))

aranjuez$year <- as.numeric(
    format(aranjuez$Date, '%Y'))

aranjuez$day <- as.numeric(
    format(aranjuez$Date, '%j'))
```

```
summary(aranjuez[, c("Date", "month", "year", "day")])
```

```
     Date                month             year          day
Min.   :2004-01-01   Min.   : 1.000   Min.   :2004   Min.   :  1.0
1st Qu.:2005-12-29   1st Qu.: 4.000   1st Qu.:2005   1st Qu.: 92.0
Median :2008-01-09   Median : 7.000   Median :2008   Median :184.0
Mean   :2008-01-03   Mean   : 6.526   Mean   :2008   Mean   :183.2
3rd Qu.:2010-01-02   3rd Qu.:10.000   3rd Qu.:2010   3rd Qu.:274.8
Max.   :2011-12-31   Max.   :12.000   Max.   :2011   Max.   :366.0
```

[1] Más información en `help(format.Date)` y `help(strptime)`.

Indexado con []

- Filas

```
aranjuez[1:5,]
```

```
        Date TempAvg TempMax TempMin HumidAvg HumidMax WindAvg WindMax Rain
1 2004-01-01   4.044   10.71  -1.969     88.3     95.9   0.746   3.528    0
2 2004-01-02   5.777   11.52   1.247     83.3     98.5   1.078   6.880    0
3 2004-01-03   5.850   13.32   0.377     75.0     94.4   0.979   6.576    0
4 2004-01-04   4.408   15.59  -2.576     82.0     97.0   0.633   3.704    0
5 2004-01-05   3.081   14.58  -2.974     83.2     97.0   0.389   2.244    0
  Radiation        ET month year day
1     5.490 0.5352688     1 2004   1
2     6.537 0.7710499     1 2004   2
3     8.810 0.8361229     1 2004   3
4     9.790 0.6861381     1 2004   4
5    10.300 0.5152422     1 2004   5
```

- Filas y Columnas

```
aranjuez[10:14, 1:5]
```

```
         Date TempAvg TempMax TempMin HumidAvg
10 2004-01-10   10.85   16.59   5.676     84.9
11 2004-01-11    7.59    9.23   4.806     95.4
12 2004-01-12    7.41   10.24   5.200     93.1
13 2004-01-13    8.35   11.38   4.137     91.3
14 2004-01-14    8.74   13.32   2.857     86.9
```

Indexado con []

- Condición basada en los datos

```
idx <- with(aranjuez, Radiation > 20 & TempAvg < 10)

head(aranjuez[idx, ])
```

```
          Date TempAvg TempMax TempMin HumidAvg HumidMax WindAvg WindMax Rain
82  2004-03-22    9.78   16.12   4.340    51.65     87.9   1.526   7.660    0
83  2004-03-23    8.50   15.52  -0.290    50.10     83.3   1.533   6.027    0
85  2004-03-25    7.47   14.58   1.584    49.66     76.6   1.138   5.939    0
100 2004-04-09    8.83   15.52   2.056    47.50     70.8   1.547   6.125    0
101 2004-04-10    7.04   13.85  -0.155    54.45     85.8   1.448   6.958    0
102 2004-04-11    7.50   15.19  -1.699    54.98     91.0   1.126   7.590    0
    Radiation       ET month year day
82      21.92 3.075785     3 2004  82
83      20.62 2.881419     3 2004  83
85      22.44 2.849603     3 2004  85
100     25.45 3.566452     4 2004 100
101     21.07 2.943239     4 2004 101
102     20.99 2.905479     4 2004 102
```

subset

```
subset(aranjuez,
       subset = (Radiation > 20 & TempAvg < 10),
       select = c(Radiation, TempAvg,
          TempMax, TempMin))
```

```
     Radiation TempAvg TempMax TempMin
82       21.92   9.780   16.12   4.340
83       20.62   8.500   15.52  -0.290
85       22.44   7.470   14.58   1.584
100      25.45   8.830   15.52   2.056
101      21.07   7.040   13.85  -0.155
102      20.99   7.500   15.19  -1.699
104      25.76   9.420   17.47   0.115
461      24.29   7.460   14.66  -0.081
462      25.25   7.930   17.35  -1.686
463      24.56   9.800   19.08  -1.484
1146     20.08   7.170   18.20  -3.746
1157     20.90   4.378   12.03  -6.353
1159     21.87   7.920   18.54  -2.941
1160     20.35   7.830   16.49  -2.807
1521     21.54   8.100   19.29  -4.075
2244     20.49   6.121   15.15  -0.940
2245     21.02   5.989   16.94  -3.208
2246     20.22   9.020   19.74  -2.068
2261     23.00   9.500   14.96   3.662
2262     20.40   9.910   14.70   4.668
2263     24.09   9.440   16.89   0.794
2265     23.64   9.680   16.35   2.938
2295     22.46   8.730   13.84   1.740
```

Ejercicio

Valores en las estaciones

Extrae dos subconjuntos de datos, uno correspondiente al invierno y otro correspondiente al verano, incluyendo las variables de radiación y temperatura media, fecha y mes.
Con estos dos `data.frame` obtén uno conjunto, diferenciando la estación de cada registro.
Puedes suponer que el invierno comenzó el 22 de diciembre y terminó el 20 de marzo, y el verano comenzó el 21 de junio y terminó el 23 de septiembre.

Solución

```
invierno <- subset(aranjuez,
                   select = c(Date, day, month,
                              Radiation, TempAvg),
                   subset = day < 79 | day > 357)
```

```
verano <- subset(aranjuez,
                 select = c(Date, day, month,
                            Radiation, TempAvg),
                   subset = day > 173 & day < 267)
```

```
invierno$id <- "Invierno"
verano$id <- "Verano"

aranjuez2 <- rbind(invierno, verano)
```

aggregate

```
aranjuez$rainy <- aranjuez$Rain > 0

aggregate(Radiation ~ rainy, data = aranjuez,
          FUN = mean)
```

```
  rainy Radiation
1 FALSE  19.63325
2  TRUE  10.26028
```

Variable categórica con cut

```
aranjuez$tempClass <- cut(aranjuez$TempAvg, 5)

aggregate(Radiation ~ tempClass, data = aranjuez,
          FUN = mean)
```

```
     tempClass Radiation
1 (-5.34,1.89]  8.805389
2  (1.89,9.09]  9.014178
3  (9.09,16.3] 14.554177
4  (16.3,23.5] 21.912414
5  (23.5,30.7] 26.192742
```

```
aggregate(Radiation ~ tempClass + rainy,
          data = aranjuez, FUN = mean)
```

```
      tempClass rainy Radiation
1  (-5.34,1.89] FALSE  9.869134
2   (1.89,9.09] FALSE 10.718837
3   (9.09,16.3] FALSE 17.238283
4   (16.3,23.5] FALSE 23.238145
5   (23.5,30.7] FALSE 26.392665
6  (-5.34,1.89]  TRUE  6.822955
7   (1.89,9.09]  TRUE  7.063932
8   (9.09,16.3]  TRUE 11.091063
9   (16.3,23.5]  TRUE 15.802522
10  (23.5,30.7]  TRUE 22.545862
```

Agregamos varias variables

```
aggregate(cbind(Radiation, TempAvg) ~ tempClass,
          data = aranjuez, FUN = mean)
```

```
     tempClass Radiation    TempAvg
1 (-5.34,1.89]  8.805389  0.3423095
2  (1.89,9.09]  9.014178  5.6663267
3  (9.09,16.3] 14.554177 12.5219084
4  (16.3,23.5] 21.912414 19.7486310
5  (23.5,30.7] 26.192742 26.0496953
```

```
aggregate(cbind(Radiation, TempAvg) ~ tempClass + rainy,
          data = aranjuez, FUN = mean)
```

```
      tempClass rainy Radiation    TempAvg
1  (-5.34,1.89] FALSE  9.869134  0.3550122
2   (1.89,9.09] FALSE 10.718837  5.6657481
3   (9.09,16.3] FALSE 17.238283 12.6959488
4   (16.3,23.5] FALSE 23.238145 19.9486604
5   (23.5,30.7] FALSE 26.392665 26.0896408
6  (-5.34,1.89]  TRUE  6.822955  0.3186364
7   (1.89,9.09]  TRUE  7.063932  5.6669887
8   (9.09,16.3]  TRUE 11.091063 12.2973563
9   (16.3,23.5]  TRUE 15.802522 18.8267565
10  (23.5,30.7]  TRUE 22.545862 25.3210345
```

Ejercicio

Valores en las estaciones

A partir del `data.frame` que incluía los datos de invierno y verano, calcula:

- La **mediana** de las variables de radiación y temperatura por estación.
- La **desviación estándar** relativa a la media de las variables de radiación y temperatura por estación.

A partir del `data.frame` completo calcula la **media** interanual diaria de las variables de radiación y temperatura.

Solución

```
## Mediana
aggregate(cbind(Radiation, TempAvg) ~ id,
          data = aranjuez2,
          FUN = median)

## Desviación estándar relativa
sdr <- function(x) sd(x) / mean(x)

aggregate(cbind(Radiation, TempAvg) ~ id,
          data = aranjuez2,
          FUN = sdr)

## Media interanual
aggregate(cbind(Radiation, TempAvg) ~ day,
          data = aranjuez,
          FUN = mean)
```

Forma simple con stack

```
aranjuezWide <- aranjuez[, c('Date', 'Radiation',
                             'TempAvg', 'TempMax',
                             'WindAvg', 'WindMax')]
```

- Pasamos de formato wide a long

```
aranjuezLong <- stack(aranjuezWide)

head(aranjuezLong)
```

```
Warning message:
In stack.data.frame(aranjuezWide) : non-vector columns will be ignored
  values       ind
1  5.490 Radiation
2  6.537 Radiation
3  8.810 Radiation
4  9.790 Radiation
5 10.300 Radiation
6  9.940 Radiation
```

```
summary(aranjuezLong)
```

```
    values              ind
Min.   :-5.309   Radiation:2898
1st Qu.: 3.158   TempAvg  :2898
Median : 8.720   TempMax  :2898
Mean   :12.074   WindAvg  :2898
3rd Qu.:19.970   WindMax  :2898
Max.   :41.910
```

Más flexible con `reshape2`

- `reshape2` es un paquete que puede facilitar la transformación de `data.frame` y matrices.

```
library(reshape2)
```

`melt` para cambiar de *wide* a *long*

```
aranjuezLong2 <- melt(aranjuezWide, id.vars = 'Date',
                     variable.name = 'Variable',
                     value.name = 'Value')

head(aranjuezLong2)
```

```
        Date  Variable  Value
1 2004-01-01 Radiation  5.490
2 2004-01-02 Radiation  6.537
3 2004-01-03 Radiation  8.810
4 2004-01-04 Radiation  9.790
5 2004-01-05 Radiation 10.300
6 2004-01-06 Radiation  9.940
```

Agregamos a partir de un formato `long`

```
aggregate(Value ~ Variable, data = aranjuezLong2,
          FUN = mean)
```

```
   Variable     Value
1 Radiation 16.741759
2   TempAvg 14.404856
3   TempMax 22.531033
4   WindAvg  1.173983
5   WindMax  5.208021
```

`dcast` para cambiar de *long* a *wide*

```
aranjuezWide2 <- dcast(aranjuezLong2,
                       Variable ~ Date)
head(aranjuezWide2[, 1:10])
```

```
Using Value as value column: use value.var to override.
   Variable 2004-01-01 2004-01-02 2004-01-03 2004-01-04 2004-01-05 2004-01-06
1 Radiation      5.490      6.537      8.810      9.790     10.300      9.940
2   TempAvg      4.044      5.777      5.850      4.408      3.081      2.304
3   TempMax     10.710     11.520     13.320     15.590     14.580     11.830
4   WindAvg      0.746      1.078      0.979      0.633      0.389      0.436
5   WindMax      3.528      6.880      6.576      3.704      2.244      2.136
  2004-01-07 2004-01-08 2004-01-09
1      7.410      4.630      4.995
2      2.080      6.405     12.060
3     11.500     13.380     15.330
4      0.449      1.188      2.737
5      3.949      6.821      7.750
```

Con merge

- Primero construimos un `data.frame` de ejemplo

```
USStates <- as.data.frame(state.x77)
USStates$Name <- rownames(USStates)
rownames(USStates) <- NULL
```

- Lo partimos en estados «fríos» y estados «grandes»

```
coldStates <- USStates[USStates$Frost>150,
                       c('Name', 'Frost')]
largeStates <- USStates[USStates$Area>1e5,
                        c('Name', 'Area')]
```

Con merge

- Unimos los dos conjuntos (estados «fríos» y «grandes»)

```
merge(coldStates, largeStates)
```

```
      Name Frost   Area
1   Alaska   152 566432
2 Colorado   166 103766
3  Montana   155 145587
4   Nevada   188 109889
```

merge usa match

- Estados grandes que también son fríos

```
idxLarge <- match(largeStates$Name,
                  coldStates$Name,
                  nomatch=0)
idxLarge
```

```
[1] 1 0 0 2 5 6 0 0
```

```
coldStates[idxLarge,]
```

```
       Name Frost
2    Alaska   152
6  Colorado   166
26  Montana   155
28   Nevada   188
```

merge usa match

- Estados frios que también son grandes

```
idxCold <- match(coldStates$Name,
                 largeStates$Name,
                 nomatch=0)
idxCold
```

```
[1] 1 4 0 0 5 6 0 0 0 0 0
```

```
largeStates[idxCold,]
```

```
       Name   Area
2    Alaska 566432
6  Colorado 103766
26  Montana 145587
28   Nevada 109889
```